

# Tomás Pablo Korenblit

tomaskorenblit@gmail.com • +541121561770 • linkedin.com/in/tomaskorenblit • tkoren.com • ORCID • Buenos Aires

I am, at heart, a maker. I was drawn to AI safety because it connects my desire to do meaningful work with technical skills. Finishing my BSc in Data Science, focused on Bayesian inference and ML. Currently doing independent research and open to collaboration.

## PUBLICATIONS

---

**Not All Instructions Are Forgotten Equal** 2026

*JAIIO 2026 (ASAIID)* • In press. Sole author [Link](#)

- Do LLMs forget some instructions faster than others? Per instruction effects span more than an order of magnitude where some can fail while mean compliance stays high
- Used a Bayesian ordered-logistic model with hierarchical effects, fit in PyMC

**PowerBench: LLM Refusal in Power-Grabbing Requests** 2026

*Apart Research Global South AI Safety Hackathon* • Coauthor, in preparation for AAAI 2027 [Link](#)

- First public benchmark for whether LLMs go along with requests to grab power from others. Varied by model, language, nationality, domain, and scale

## EXPERIENCE

---

**AI Safety** • Independent Researcher Apr 2026 – Present

- Working with a team on a multilingual benchmark for LLM refusal in power-grabbing requests, for submission to AAAI 2027
- Participating in BlueDot Impact projects, self-studying ARENA

**Ascendancy** • Partner & Founding Data Scientist Jan 2026 – Apr 2026

- Built a graph of investors, founders, and companies across the Americas, MENA, and Europe, with warm-intro pathfinding and community detection over it
- Built an assistant that turns plain English into Cypher, so non-technical users can query the graph directly
- Python • Neo4j • PostgreSQL • FastAPI • GCP

**Invisible Technologies** Dec 2022 – Dec 2025

*Software Engineer* (Jun 2025) • *QC Data Analyst* (Jun 2023) • *AI Data Trainer* (Dec 2022)

- Built a RAG platform for creating RLHF annotation interfaces, used across the trainer teams
- Built validation tools and dashboards (Python, Looker) used across the team. Reviewed output quality for the LLM training teams

**Klouser** • Data Scientist / Data Engineer Dec 2021 – Oct 2022

- Built a prospect-generation system and trained a classifier on client profiles to qualify leads (Python, Pandas, PostgreSQL, Scikit-learn)

## EDUCATION

---

**BSc Data Science** • Universidad Nacional de San Martín 2021 – 2026 (expected)

**Technical AI Safety** • BlueDot Impact 2026

## SKILLS

---

Bayesian & ML: Bayesian Stats, PyMC, PyTorch, Scikit-learn

LLM evaluation: Inspect AI, RAG, RLHF tooling

Languages: Python, R, Julia, SQL, Cypher, TypeScript

Data & Backend: Neo4j, PostgreSQL, Redis, Pandas, NumPy | FastAPI, Flask, Pydantic, SQLAlchemy

Infrastructure: GCP, Docker

Spoken: English & Spanish