

Tomás Pablo Korenblit

EDUCATION	Universidad Nacional de San Martín BSc Data Science	2021 – 2026 (expected)
PUBLICATIONS	Not All Instructions Are Forgotten Equal <i>JAIIO 2026 (ASAIID)</i> . In press. Sole author	2026 paper repo
	– Aggregate compliance hides per-instruction failures. Per-instruction treatment effects span more than an order of magnitude, so a model can drop a stated preference while mean compliance stays high – Stating a preference a model already follows can backfire. Told not to add new imports, models added more. Explicit negation seems to push the model off a good default – Hierarchical Bayesian ordered-logistic model fit in PyMC. Scored with deterministic checkers, no LLM in the loop	
	PowerBench: LLM Refusal in Power-Grabbing Requests <i>Apart Research Global South AI Safety Hackathon</i> . Coauthor. In preparation	2026 paper
	– Adding self-benefit to a harmful request lowers refusal. Models refuse pure disempowerment (harm to others) more than power-grabbing (harm plus self-gain). Holds across five models – Refusal of these requests does not correlate with model capability, so it reflects training choices – Multilingual, single-turn first-person requests, scored by an LLM judge not yet validated against human labels	
EXPERIENCE	Research – Studying LLM threat models. Extending the PowerBench work – Facilitating BlueDot’s Technical AI Safety course, and taking their Projects course	Apr 2026 – Present
	Ascendancy <i>Data Scientist</i> – Built a graph of investors, founders, and companies across the Americas, MENA, and Europe, with warm-intro pathfinding and community detection over it – Wrote an assistant to turn English into Cypher, so non-technical users can query the graph directly (Python, Neo4j, PostgreSQL, FastAPI, GCP)	Jan 2026 – Apr 2026
	Invisible Technologies <i>Software Engineer</i> (Jun 2025) <i>QC Data Analyst</i> (Jun 2023) <i>AI Data Trainer</i> (Dec 2022) – Designed a RAG platform to create almost all RLHF annotation interfaces – Wrote validation tools and dashboards (Python, Looker) used company-wide. Reviewed output quality for the LLM training teams	Dec 2022 – Dec 2025
	Klouser <i>Data Scientist</i> / <i>Data Engineer</i> – Made a lead generation system and trained a classifier on client profiles to qualify them (Python, Pandas, PostgreSQL, Scikit-learn)	Dec 2021 – Oct 2022